

Analyzing Categorical Data

In this project, we'll be working with categorical data and will be using a subset of data from the following data set: (<https://www.kaggle.com/datasets/norc/general-social-survey?select=gss.csv>) (<https://www.kaggle.com/datasets/norc/general-social-survey?select=gss.csv>)).

After cleaning the data, we will use some visualizations tools. We also had used statsmodels for a special type of categorical plot.

```
In [3]: # Import packages
import pandas as pd
import numpy as np
import plotly.express as px
import matplotlib.pyplot as plt
from statsmodels.graphics.mosaicplot import mosaic

# Read in csv as a DataFrame and preview it
df = pd.read_csv("/Users/antoniogondim/Downloads/gss_sub.csv")
df
```

Out[3]:

	year	id	labor_status	self_employed	marital_status	n_siblings	age	high_school	
0	1972.0	1.0	WORKING FULLTIME	SOMEONE ELSE	NEVER MARRIED	3.0	23.0	16.0	BA
1	1972.0	2.0	RETIRED	SOMEONE ELSE	MARRIED	4.0	70.0	10.0	ξ
2	1972.0	3.0	WORKING PARTTIME	SOMEONE ELSE	MARRIED	5.0	48.0	12.0	ξ
3	1972.0	4.0	WORKING FULLTIME	SOMEONE ELSE	MARRIED	5.0	27.0	17.0	BA
4	1972.0	5.0	KEEPING HOUSE	SOMEONE ELSE	MARRIED	2.0	61.0	12.0	ξ
...	
59594	2014.0	2539.0	KEEPING HOUSE	SOMEONE ELSE	WIDOWED	6.0	89.0	14.0	C
59595	2014.0	2540.0	WORKING FULLTIME	SOMEONE ELSE	DIVORCED	3.0	56.0	12.0	ξ
59596	2014.0	2541.0	WORKING FULLTIME	SOMEONE ELSE	NEVER MARRIED	5.0	24.0	14.0	ξ
59597	2014.0	2542.0	WORKING FULLTIME	SOMEONE ELSE	NEVER MARRIED	2.0	27.0	13.0	ξ
59598	2014.0	2543.0	WORKING PARTTIME	SOMEONE ELSE	WIDOWED	2.0	71.0	12.0	ξ

59599 rows × 16 columns

In [4]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 59599 entries, 0 to 59598
Data columns (total 16 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   year                                59599 non-null  float64
1   id                                  59599 non-null  float64
2   labor_status                       59583 non-null  object
3   self_employed                     59306 non-null  object
4   marital_status                     59575 non-null  object
5   n_siblings                         56682 non-null  float64
6   age                                59599 non-null  float64
7   high_school                        59440 non-null  float64
8   degree                             59464 non-null  object
9   political_affiliation              59257 non-null  object
10  environment                         59388 non-null  object
11  law_enforcement                    59378 non-null  object
12  drugs                              59380 non-null  object
13  space_exploration                  59596 non-null  object
14  inequality                          1532 non-null   float64
15  household_size                     59599 non-null  float64
dtypes: float64(7), object(9)
memory usage: 7.3+ MB
```

In [5]:

```
df=df.drop('inequality', axis=1)
#Too many null values at this column
df.isnull().sum()
```

```
Out[5]: year                                0
id                                  0
labor_status                       16
self_employed                     293
marital_status                     24
n_siblings                         2917
age                                0
high_school                        159
degree                             135
political_affiliation              342
environment                         211
law_enforcement                    221
drugs                              219
space_exploration                   3
household_size                     0
dtype: int64
```

```
In [6]: df['n_siblings'].value_counts()
```

```
Out[6]: 2.0      10717
        1.0      9602
        3.0      9109
        4.0      6705
        5.0      4860
        6.0      3688
        7.0      3168
        8.0      2075
        9.0      1536
       -1.0      1518
       10.0     1066
       11.0      826
       12.0      550
       13.0      370
       14.0      213
       15.0      144
       99.0      111
       16.0      105
       98.0       57
       17.0       48
       21.0       46
       18.0       41
       20.0       29
       19.0       25
       22.0       15
       23.0       14
       25.0        6
       27.0        6
       26.0        6
       31.0        6
       24.0        5
       30.0        4
       33.0        2
       32.0        2
       29.0        2
       68.0        1
       35.0        1
       34.0        1
       37.0        1
       55.0        1
      Name: n_siblings, dtype: int64
```

```
In [7]: df.n_siblings.replace(np.nan,int(2.0), inplace=True)
```

```
In [8]: df.shape
```

```
Out[8]: (59599, 15)
```

Type *Markdown* and LaTeX: α^2

```
In [9]: df=df.dropna()
df.shape
```

```
Out[9]: (58439, 15)
```

Above we see that our DataFrame contains `float64` column (numerical data), as well as a number of `object` columns, i.e object data types contain strings.

`df.describe()` method with the `include` parameter to select a particular DataType (in this case "o"). This returns the count, number of unique values, the mode, and frequency of the mode for each column having object as data type.

```
In [10]: df.describe(include="o")
```

```
Out[10]:
```

	labor_status	self_employed	marital_status	degree	political_affiliation	environment	law_ei
count	58439	58439	58439	58439	58439	58439	
unique	8	4	5	6	9	5	
top	WORKING FULLTIME	SOMEONE ELSE	MARRIED	HIGH SCHOOL	NOT STR DEMOCRAT	IAP	
freq	28960	48809	31376	30124	12272	37576	

```
In [11]: df["environment"].value_counts()
```

```
Out[11]: IAP          37576
TOO LITTLE    12971
ABOUT RIGHT   5351
TOO MUCH      1660
DK             881
Name: environment, dtype: int64
```

Manipulating categorical data

- The categorical variable type can be useful, especially here:
 - It is possible to specify a precise order to the categories when the default order may be incorrect (e.g., via alphabetical).
 - Can be compatible with other Python libraries.

Let's take our existing categorical variables and convert them from strings to categories. Here, we use `.select_dtypes()`.

(https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.select_dtypes.html) to return only object columns, and with a dictionary set their type to be a category.

```
In [13]: # Create a dictionary of column and data type mappings
conversion_dict = {k: "category" for k in df.select_dtypes(include="object")}

# Convert our DataFrame and check the data types
df = df.astype(conversion_dict)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 58439 entries, 0 to 59598
Data columns (total 15 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   year                  58439 non-null  float64
 1   id                    58439 non-null  float64
 2   labor_status          58439 non-null  category
 3   self_employed         58439 non-null  category
 4   marital_status        58439 non-null  category
 5   n_siblings            58439 non-null  float64
 6   age                   58439 non-null  float64
 7   high_school           58439 non-null  float64
 8   degree                58439 non-null  category
 9   political_affiliation 58439 non-null  category
10   environment           58439 non-null  category
11   law_enforcement       58439 non-null  category
12   drugs                 58439 non-null  category
13   space_exploration     58439 non-null  category
14   household_size        58439 non-null  float64
dtypes: category(9), float64(6)
memory usage: 3.6 MB
```

Already we can see that the memory usage of the DataFrame has been halved from 7 mb to about 4 mb, optimizing the data.

Cleaning up the labor_status column

To analyze the relationship between employment and attitudes over time, we need to clean up the labor_status column. We can preview the existing categories using .categories .

```
In [14]: df["labor_status"].values
```

```
Out[14]: ['WORKING FULLTIME', 'RETIRED', 'WORKING PARTTIME', 'WORKING FULLTIME',
'KEEPING HOUSE', ..., 'KEEPING HOUSE', 'WORKING FULLTIME', 'WORKING FULLT
IME', 'WORKING FULLTIME', 'WORKING PARTTIME']
Length: 58439
Categories (8, object): ['KEEPING HOUSE', 'OTHER', 'RETIRED', 'SCHOOL',
'TEMP NOT WORKING', 'UNEMPL, LAID OFF', 'WORKING FULLTIME', 'WORKING PART
TIME']
```

```
In [15]: df["labor_status"].value_counts()
```

```
Out[15]: WORKING FULLTIME      28960
KEEPING HOUSE      9478
RETIRED           7861
WORKING PARTTIME   6012
UNEMPL, LAID OFF   1920
SCHOOL            1807
TEMP NOT WORKING   1240
OTHER             1161
Name: labor_status, dtype: int64
```

Let's collapse some of these categories. The easiest way to do this is to replace the values inside the column using a dictionary, and then reset the data type back to a category.

```
In [16]: # Create a dictionary of categories to collapse
new_labor_status = {"UNEMPL, LAID OFF": "UNEMPLOYED",
                    "TEMP NOT WORKING": "UNEMPLOYED",
                    "WORKING FULLTIME": "EMPLOYED",
                    "WORKING PARTTIME": "EMPLOYED"
                  }

# Replace the values in the column and reset as a category
df["labor_status_clean"] = df["labor_status"].replace(new_labor_status).astype('category')
print(df.dtypes)
# Preview the new column
df["labor_status_clean"].value_counts()
```

```
year          float64
id            float64
labor_status   category
self_employed  category
marital_status category
n_siblings     float64
age            float64
high_school    float64
degree         category
political_affiliation category
environment     category
law_enforcement category
drugs           category
space_exploration category
household_size float64
labor_status_clean category
dtype: object
```

```
Out[16]: EMPLOYED      34972
KEEPING HOUSE      9478
RETIRED           7861
UNEMPLOYED        3160
SCHOOL            1807
OTHER             1161
Name: labor_status_clean, dtype: int64
```

Reordering categories

```
In [17]: df["environment"].values
```

```
Out[17]: ['IAP', 'IAP', 'IAP', 'IAP', 'IAP', ..., 'TOO LITTLE', 'TOO LITTLE', 'TOO LITTLE', 'IAP', 'IAP']
Length: 58439
Categories (5, object): ['ABOUT RIGHT', 'DK', 'IAP', 'TOO LITTLE', 'TOO MUCH']
```

```
In [21]: # Set the new order
new_order = ["TOO LITTLE", "ABOUT RIGHT", "TOO MUCH", "DK", "IAP"]
categories_to_remove = ["DK", "IAP"]

# Loop through each column
for col in ["environment", "law_enforcement", "drugs"]:
    # Reorder and remove the categories
    df[col + "_clean"] = df[col].cat.reorder_categories(new_order, ordered=True)
    df[col + "_clean"] = df[col + "_clean"].cat.remove_categories(categories_to_remove)

# Preview one of the columns' categories
df["environment_clean"].cat.categories
```

```
Out[21]: Index(['TOO LITTLE', 'ABOUT RIGHT', 'TOO MUCH'], dtype='object')
```

Now let's also apply these steps to education level in one go: collapsing, removing, and reordering.

```
In [22]: df['degree'].values #let's reorder that and remove 'DK'
```

```
Out[22]: ['BACHELOR', 'LT HIGH SCHOOL', 'HIGH SCHOOL', 'BACHELOR', 'HIGH SCHOOL', ..., 'JUNIOR COLLEGE', 'HIGH SCHOOL', 'HIGH SCHOOL', 'HIGH SCHOOL', 'HIGH SCHOOL']
Length: 58439
Categories (6, object): ['BACHELOR', 'DK', 'GRADUATE', 'HIGH SCHOOL', 'JUNIOR COLLEGE', 'LT HIGH SCHOOL']
```

```

In [23]: # Define a dictionary to map old degree categories to new ones
new_degree = {"LT HIGH SCHOOL": "HIGH SCHOOL",
              "BACHELOR": "COLLEGE/UNIVERSITY",
              "GRADUATE": "COLLEGE/UNIVERSITY",
              "JUNIOR COLLEGE": "COLLEGE/UNIVERSITY"}

# Replace old degree categories with new ones and convert to categorical data
df["degree_clean"] = df["degree"].replace(new_degree).astype("category")

# Remove "DK" category from degree_clean column
df["degree_clean"] = df["degree_clean"].cat.remove_categories(["DK"])

# Reorder degree_clean categories and set as ordered
df["degree_clean"] = df["degree_clean"].cat.reorder_categories(["HIGH SCHOOL", "COLLEGE/UNIVERSITY"])

# Preview the new column
df["degree_clean"].value_counts()

```

```

Out[23]: HIGH SCHOOL          42756
         COLLEGE/UNIVERSITY    15660
         Name: degree_clean, dtype: int64

```

By [IntervalIndex](https://pandas.pydata.org/docs/reference/api/pandas.IntervalIndex.html) (<https://pandas.pydata.org/docs/reference/api/pandas.IntervalIndex.html>) we set cutoff ranges for the `year` . We then use `pd.cut()` (<https://pandas.pydata.org/docs/reference/api/pandas.cut.html>) to cut our `year` column by these ranges, and set labels for each range.


```
In [27]: decade_boundaries = [(1970, 1979), (1979, 1989), (1989, 1999), (1999, 2009),
# Set the bins and cut the DataFrame

bins = pd.IntervalIndex.from_tuples(decade_boundaries) #this method returns

decade_labels = {bins[0]: '1970s',
                  bins[1]: '1980s',
                  bins[2]: '1990s',
                  bins[3]: '2000s',
                  bins[4]: '2010s'}

print(bins)
df["decade"] = pd.cut(df["year"], bins) #creates a new column based on the y
print(df['decade'].values, df['decade'].dtypes)

# Rename each of the intervals of decade_boundaries as the decades on decad
df['decade'] = df["decade"].replace(decade_labels) #.astype('category')

# Preview the new column
df[['year', 'decade']]
```

```
IntervalIndex([(1970, 1979], (1979, 1989], (1989, 1999], (1999, 2009], (2009, 2019]], dtype='interval[int64, right]')
[(1970, 1979], (1970, 1979], (1970, 1979], (1970, 1979], (1970, 1979],
..., (2009, 2019], (2009, 2019], (2009, 2019], (2009, 2019], (2009, 2019]]
Length: 58439
Categories (5, interval[int64, right]): [(1970, 1979] < (1979, 1989] < (1989, 1999] < (1999, 2009] < (2009, 2019]] category
```

Out[27]:


	year	decade
0	1972.0	1970s
1	1972.0	1970s
2	1972.0	1970s
3	1972.0	1970s
4	1972.0	1970s
...
59594	2014.0	2010s
59595	2014.0	2010s
59596	2014.0	2010s
59597	2014.0	2010s
59598	2014.0	2010s

58439 rows × 2 columns

Visualization

```
In [28]: # Create a new figure object
fig = px.bar(df["labor_status"].value_counts(),
             template="plotly_white")

# Hide the legend and show the plot
fig.update_layout(showlegend=False)
fig.show()
```




Let's change the orientation of the plot and add a title, for a better perspective.

```
In [96]: # Create a new figure object
fig = px.bar(df["labor_status"].value_counts(ascending=True),
             template="plotly_white",
             orientation="h",
             title="Labor status by count"
            )

# Hide the legend and show the plot
fig.update_layout(showlegend=False)
fig.show()
```

Labor status by count



Bar charts

```
In [97]: ## Aggregate household size by year  
household_by_decade = df.groupby("decade", as_index=False)["household_size"]  
household_by_decade
```


Out[97]:

	decade	household_size
0	1970s	3.238023
1	1980s	2.741394
2	1990s	2.546036
3	2000s	2.464682
4	2010s	2.411328

```
In [98]: # Create a new figure object
fig = px.bar(household_by_decade,
              x="decade",
              y="household_size",
              template="plotly_white",
              title="Average household size by decade"
            )

fig.show()
```

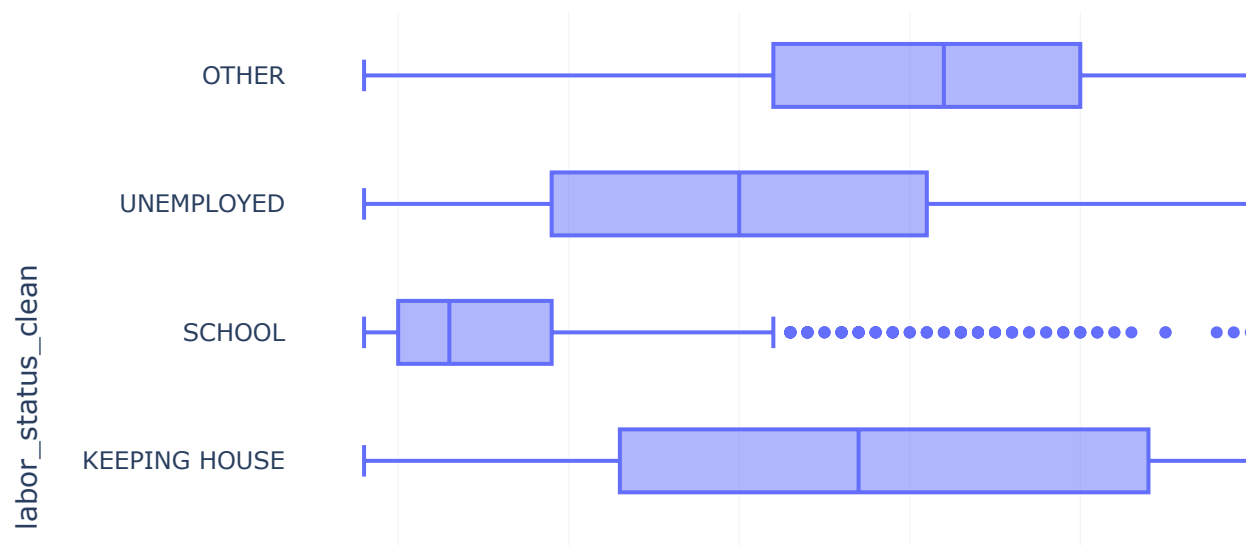
Average household size by decade



Boxplots

```
In [99]: # Create a new figure object
fig = px.box(df,
             x="age",
             y="labor_status_clean",
             template="plotly_white"
            )

fig.show()
```



Mosaic plots

visualize the relationship between two categorical variables. One way to do this is a frequency table, which will give the counts across the different combinations of the two variables. create a frequency table using `pd.crosstab()`.

(<https://pandas.pydata.org/docs/reference/api/pandas.crosstab.html>)


```
In [100]: pd.crosstab(df["degree_clean"], df["law_enforcement_clean"])
```

```
Out[100]:
```

law_enforcement_clean	TOO LITTLE	ABOUT RIGHT	TOO MUCH
degree_clean			
HIGH SCHOOL	7937	4799	1309
COLLEGE/UNIVERSITY	3193	2598	422

```
In [101]: # Create a mosaic plot and show it
mosaic(df,
        ['degree_clean', 'law_enforcement_clean'],
        title='Law enforcement opinions by degree')

plt.show()
```



Line charts

The final plot type we will cover is a line plot. Line plots often (but not always!) show the relationship between time and a numerical variable. Adding in a categorical variable can be a great way to enrich a line plot and provide other information.

Here, we use the `.value_counts()` method as an aggregation function, and use this in combination with a Plotly `line_plot()` (<https://plotly.com/python/line-charts/>) to visualize the trend in marital statuses over the years.

```
In [102]: # Group the dataframe by year and marital status, and calculate the normalized
marital_rates = df.groupby(["year"], as_index=False)["marital_status"].value_counts()

# Display the resulting DataFrame
marital_rates
```

Out[102]:

	year	marital_status	proportion
0	1972.0	MARRIED	0.723391
1	1972.0	NEVER MARRIED	0.128744
2	1972.0	WIDOWED	0.084130
3	1972.0	DIVORCED	0.039516
4	1972.0	SEPARATED	0.024219
...
145	2014.0	MARRIED	0.458918
146	2014.0	NEVER MARRIED	0.265331
147	2014.0	DIVORCED	0.163126
148	2014.0	WIDOWED	0.081363
149	2014.0	SEPARATED	0.031263


150 rows × 3 columns


```
In [103]: # Create a new figure object
fig = px.line(marital_rates,
              x="year",
              y="proportion",
              color="marital_status",
              template="plotly_white",
              title="Marital status over time"
            )

# Update the y-axis to show percentages
fig.update_yaxes(tickformat=".0%")

# Show the plot
fig.show()
```

Marital status over time



In []:

In []:

